

Sparse Planning Graphs for Information Driven Exploration

1,2,3

The Robotics Institute
Carnegie Mellon University
Pittsburgh, PA 15217
Email: {1,2,3}@cmu.edu

I. INTRODUCTION

Exploration is a key capability that enables robotic vehicles to operate in unknown environments. In general, the shortest trajectory over which a robot can greedily explore an environment is the trajectory formed by choosing actions which maximally reduce uncertainty in the environmental belief representation, or map. One strategy to accomplish this is to choose control actions that maximize an information metric between the robot's current map and the robot's map at a future timestep. This solution is one variant of a broad category of exploration strategies known as active Simultaneous Localization and Mapping (SLAM) [?].

State of the art active SLAM implementations generally only compute actions over a one-step planning horizon due to the extremely high computational cost of determining expected information gain over all potential future locations [?], [?]. However, plans over much longer horizons can be generated if expected information gain values can be computed once per planning step, cached, and updated efficiently based on new information. An even more efficient approach would utilize sparse planning graphs to limit expected information gain calculations to a select few feasible future locations in the map. Rapidly Exploring Random Trees (RRT) and lattice graphs are two examples of sparse planning graphs which decompose the reachable space into a sparse set of goal states based on a set of motion primitives for planning purposes.

This project seeks to develop a recursive formulation for efficiently and intelligently updating expected information gain over a finite horizon as the robot moves through an unknown environment and builds a map from sensor observations.

II. OCCUPANCY GRID MAPPING

As a basis for the core formulation in the following sections, we provide a brief overview of occupancy grid mapping. Occupancy grids are a common and useful We represent the map as an occupancy grid, which consists of a set of cells: $m = \{m^i\}_{i=1}^N$. The probability that an individual cell is occupied is given by $p(m^i | x_{1:t}, z_{1:t})$, where $x_{1:t}$ denotes the history of states of the vehicle, and $z_{1:t}$ denotes the history of range observations accumulated by the

vehicle. We assume that cell occupancy probabilities are independent of one another: $p(m | x_{1:t}, z_{1:t}) = \prod_i p(m^i | x_{1:t}, z_{1:t})$. For notational simplicity we write the map conditioned on random variables $x_{1:t}$ and $z_{1:t}$ as $p_t(m) := p(m | x_{1:t}, z_{1:t})$. Additionally, unobserved grid cells are assigned a uniform prior of being occupied.

We represent the occupancy status of grid cell m^i at time t with a log odds expression

$$l_t := \log \frac{p(m^i | z_{1:t})}{p(\bar{m}^i | z_{1:t})} \quad (1)$$

where \bar{m}^i denotes the probability that m^i is unoccupied. When a new observation z_t is obtained, the log odds update is given by

$$l_t = l_{t-1} + \log \frac{p(m^i | z_t)}{p(m^i)} - \log \frac{p(\bar{m}^i | z_t)}{p(\bar{m}^i)} \quad (2)$$

where the last two terms represent the inverse sensor model.

III. EXPLORATIVE INFORMATION COST FUNCTION

In this section, we will introduce an information-theoretic formulation for planning explorative paths. This formulation does not retain or propagate information from previous iterations, and is therefore slow.

The goal of exploration is to find a dynamically feasible sequential set of motions chosen from a library of motion primitives, \mathcal{X} , over a time interval, $\tau := t + 1 : t + T$, which enable the robot to explore its environment. We use the criteria that an explorative action is one which allows the robot to position itself in locations that generate observations which are informative to the robot's map. By this criteria, choosing the optimal exploration action will maximize an information metric over the robot's *future map*. In this context, a *future map* is an updated version of the robot's current map, which has used Eq. (2) to incorporate all measurements gathered while executing an action. An *action* can be defined as a discrete sequence of states, $x_\tau = [x_{t+1}, \dots, x_{t+T}]$. While executing an action, the robot will obtain a set of measurements $z_\tau(x_\tau) = [z_{t+1}(x_{t+1}), \dots, z_{t+T}(x_{t+T})]$ by sensing from the states x_τ . $z_\tau(x_\tau)$ is modeled as a random

variable whose distribution is parameterized by a deterministic action, x_τ , generated by a planner. Under this notation, an explorative planner must determine x_τ^* : the action that visits locations which allow the robot to obtain the set of measurements which are most informative to the current map. When integrated with Eq. (2), these measurements will maximize an information-theoretic cost function over the map. We choose to maximize Shannon Mutual Information (MI) rate between the current map, m , and the measurements z_τ gathered along x_τ .

$$x_\tau^* = \underset{x_\tau \in \mathcal{X}^T}{\operatorname{argmax}} \frac{I_{\text{MI}}[m; z_\tau(x_\tau)]}{R(x_\tau)} = \underset{x_\tau \in \mathcal{X}^T}{\operatorname{argmax}} \frac{H[m] - H[m | z_\tau(x_\tau)]}{R(x_\tau)} \quad (3)$$

where I_{MI} is the Shannon Mutual Information, and $R: \mathcal{X}^T \rightarrow \mathbb{R}^+$ is a function which returns the time required to execute an action. In contrast to MI, MI rate is used so that actions with different execution times can be compared under a common metric.

Fortunately, $H[m]$ depends only on the current map, and can be computed once for all enumerations of x_τ under consideration. Several additional assumptions can be made to simplify the optimization of x_τ^* . Due to cell independence in the occupancy grid formulation, the joint conditional entropy over the map can be expressed as a sum of individual cell conditional entropies. Additionally, let \mathcal{C} be the set of cells in the map that beams in z_τ pass through. Given the inverse sensor model (shown in the last two terms of Eq. (2)), cells $c \notin \mathcal{C}$ are not updated by z_τ and therefore do not contribute to the map's conditional entropy. Using these assumptions, we may write the entropy of the map conditioned on z_τ as

$$\begin{aligned} H[m | z_\tau] &= \sum_{c \in \mathcal{C}} H[c | z_\tau] \\ &= \sum_{c \in \mathcal{C}} \mathbb{E}_{c, z_\tau} [-\log p(c | z_\tau)] \\ &= - \int_{z_\tau} p(z_\tau) \sum_{c \in \mathcal{C}} o(c | z_\tau) dz_\tau \end{aligned} \quad (4)$$

where

$$o(c | z_\tau) = p(c | z_\tau) \log p(c | z_\tau) + (1 - p(c | z_\tau)) \log(1 - p(c | z_\tau)) \quad (5)$$

Computing this term requires integration over the space of measurements, z_τ , which is intractable for an online planner. Instead, we approximate Eq. (4) with N Monte Carlo samples, $\{z_\tau^i\}_{i=1}^N$, drawn from the distribution $p(z_\tau)$.

$$H[m | z_\tau] \approx -\frac{1}{N} \sum_{i=1}^N \sum_{c \in \mathcal{C}} o(c | z_\tau^i) \quad (6)$$

Substituting Eq. (6) into Eq. (3) yields an information-theoretic cost function, which, when

maximized, computes an explorative action.

$$x_\tau^* = \underset{x_\tau \in \mathcal{X}^T}{\operatorname{argmax}} \frac{1}{R(x_\tau)} \left(H[m] + \sum_{i=1}^N \sum_{c \in \mathcal{C}} o(c | z_\tau^i(x_\tau)) \right) \quad (7)$$

This cost function depends on the ability to sample from $p(z_\tau)$, necessitating a reliable measurement model from which to generate representative measurements.

IV. MEASUREMENT MODEL

As explained in Sect. 4, each future position x_j will generate a multi-beam range measurement z_j . We express each measurement as a K -tuple random variable, $z_j = [z_{j,1}, \dots, z_{j,K}]$, with $z_{j,k} \in [z_{\min}, z_{\max}]$. A measurement model is required to sample from the distribution $p(z_\tau)$. Assuming laser beam independence, $p(z_\tau)$ can be expressed as a product of the measurement probabilities of individual beams at each timestep over the interval $\tau = t+1 : t+T$.

$$p(z_\tau) = \prod_{j=t+1}^{t+T} \prod_{k=1}^K p(z_{j,k}) \quad (8)$$

The probability of a single beam measurement can be expressed as a function of the true distance, r , to the obstacle.

$$p(z_{j,k}) = \int_r p(z_{j,k} | r) p(r) dr \quad (9)$$

$$p(z_{j,k} | r) = \begin{cases} \mathcal{N}(z_{j,k} - r, \sigma_{\text{hit}}^2) & : z_{\min} \leq r \leq z_{\max} \\ \mathcal{N}(z_{\max}, \sigma_{\text{hit}}^2) & : z_{\max} < r \\ 0 & : r < z_{\min} \end{cases} \quad (10)$$

where $\mathcal{N}(x - \mu, \sigma^2)$ is a one-dimensional Gaussian with mean μ and variance σ^2 . The Gaussian term in Eq. (10) is an approximation for the true distribution, which is a complex function of both beam angle as well as non-axial distance-dependent noise. The environment-dependent distribution $p(r)$ corresponds to the probability that an obstacle intersects a beam at range r . We model this as a Poisson binomial distribution with probabilities decaying according to the inverse square of the beam range.

On the interval $[z_{\min}, z_{\max}]$, $p(z_{j,k})$ reduces to a Gaussian convolution with $p(r)$ over r , which can be computed in $O(n \log n)$ operations with a Fast Fourier Transform. This computation can be reduced to $O(n)$ by assuming that σ_{hit} of the range sensor is much smaller than the resolution of the map.

The distribution $p(r)$ represents the probability that the first obstacle along a ray originating at the sensor is found at distance r . $p(r)$ can be expressed as a generalization of the geometric distribution for independent, but not identically distributed (i.n.i.d.)




Fig. 1: The distribution $p(r | m)$ over a depicted 1-dimensional map with resolution $\Delta r = 0.1$ m, and $p_e = 0.01$, $p_o = 0.92$, $p_u = 0.05$.




Fig. 2: Snapshots of an RRT exploring the environment defined by the occupancy grid.

are currently direct connections between nodes, but there also exist variants of the algorithm that if they simulate the vehicle dynamics as part of the expansion process [?]. This approach is traditionally used to ensure dynamic feasibility and dense collision checking. However, this also provides a means for evaluating the information gain along each candidate trajectory. As a result, we can augment the cost of each branch with a metric based on the information cost function from Sect. 4. This will enable the planner to select the path in the tree that maximizes the predicted information gain.

Bernoulli trials.

$$\begin{aligned}
 p(r = i) &= p_t(m^i) \prod_{j=1}^{i-1} (1 - p_t(m^j)) \\
 &= p_t(m^i) \left(\frac{1}{p_t(m^{i-1})} - 1 \right) p(r = i - 1)
 \end{aligned} \tag{11}$$

A. Planning with RRT

To make use of this information cost function to guide exploration, we consider a sampling-based planning approach that can evaluate the predicted information gain throughout the environment. In addition, we wish to use the occupancy grid that is being updated online (as described in Sect. 3) to guide the vehicle around obstacles in the environment. The rapidly-exploring random tree (RRT) algorithm is well suited to planning paths through these types of large environments, and works as follows. The planner starts from the vehicle's current state and samples a point x in the environment. Using the occupancy grid, we can reject samples that lie in cells with a sufficiently high probability of being occupied. If the sample is valid, we find the closest node in the tree of paths (initially just the vehicle state), where closeness is measured in terms of Euclidean distance, and add a new edge to the tree connecting the sample point to the nearest node. Then a new sample is drawn and the process repeats to grow a tree of path segments through the environment. This tree-growing process terminates after a specified time, and the minimum cost path is returned.

Figure 2 shows snapshots of our current implementation of RRT planning over the occupancy grid. As seen in the figure, the edges in the tree

V. UNSCENTED KALMAN FILTER